



School of Computer Science and Engineering  
Faculty of Engineering  
The University of New South Wales

# **A Modern IEEE 2030.5 Client Implementation**

by Ethan Dickson

Supervised by Jawad Ahmed  
Assessed by Nadeem Ahmed

**Thesis A Report**  
Submitted April 2023

Thesis submitted as a requirement for the degree of  
Bachelor of Engineering in Software Engineering

# Abstract

IEEE 2030.5 is a standardised application-layer communications protocol between end-user energy devices, and electric utility management systems. In this thesis, we'll be examining predecessors to the protocol, competing protocols, and existing implementations of the standard. This thesis will then analyze how modern programming techniques can be used to produce a safe, secure, robust, reliable, and open-source implementation of an IEEE 2030.5 client library for use in end-user energy devices.

# Acknowledgements

I would like to thank Jawad Ahmed and Nadeem Ahmed for their guidance as supervisor and assessor, respectively.

I would also like to thank Neel Bhaskar and Jack Peng, for their work on an IEEE 2030.5 server implementation as part of research at UNSW previously.

Finally, I would like to thank my friends and family for their support during my degree and thesis.

# Abbreviations

**API** Application Programming Interface

**ARENA** Australian Renewable Energy Agency

**CPUC** California Public Utilities Commission

**CPU** Central Processing Unit

**CSE** Computer Science and Engineering

**CVE** Common Vulnerabilities and Exposures

**DER** Distributed Energy Resources

**DNS** Domain Name System

**DNSP** Distribution Network Service Provider

**DNS-SD** Domain Name System - Service Discovery

**ECC** Electric Curve Cryptography

**EPRI** Electric Power Research Institute

**EXI** Efficient XML Interchange

**FS** Function Set

**HAN** Home Area Network

**HTTP** Hypertext Transfer Protocol

**IEEE** Institute of Electrical and Electronics Engineers

**I/O** Input / Output

**IoT** Internet of Things

**IP** Internet Protocol

**OOP** Object-Oriented Programming

**OS** Operating System

**REST** Representational State Transfer

**RSA** Rivest-Shamir-Adleman

**SEP** Smart Energy Profile

**SIP** Smart Inverter Profile

**TCP** Transmission Control Protocol

**TLS** Transport Layer Security

**UDP** User Datagram Protocol

**UNSW** University of New South Wales

**XML** Extensible Markup Language

**XSD** XML Schema Definition

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Context</b>	<b>2</b>
2.1	Smart Energy Profile 1.x . . . . .	2
2.2	SunSpec Modbus . . . . .	2
2.3	IEEE 2030 . . . . .	3
2.4	Producing IEEE 2030.5 . . . . .	3
<b>3</b>	<b>Background</b>	<b>4</b>
3.1	High-Level Architecture . . . . .	4
3.2	Protocol Design . . . . .	5
3.3	Usage . . . . .	5
<b>4</b>	<b>Adoption</b>	<b>6</b>
4.1	California Public Utilities Commission . . . . .	6
4.2	Australian Renewable Energy Agency . . . . .	8
4.3	SunSpec Alliance . . . . .	9
4.4	Electric Power Research Institute . . . . .	9
<b>5</b>	<b>Implementation Design</b>	<b>10</b>
5.1	Programming Language . . . . .	10
5.1.1	Rust . . . . .	11
5.2	Operating System . . . . .	12
5.3	Common Library . . . . .	12
5.3.1	Resource Data Types . . . . .	12
5.3.2	Resource Serialisation & Deserialisation . . . . .	17
5.3.3	Security . . . . .	18
5.3.4	Network . . . . .	18
5.3.5	Testing . . . . .	18
5.4	Client Library . . . . .	20
5.4.1	Server Discovery & Connection . . . . .	20
5.4.2	Event-driven Architecture . . . . .	21
5.4.3	Function Sets . . . . .	21
5.4.4	Testing . . . . .	23
<b>6</b>	<b>Timeline</b>	<b>24</b>
6.1	Thesis A . . . . .	24
6.2	Thesis B . . . . .	25
6.3	Thesis C . . . . .	25
	<b>Bibliography</b>	<b>27</b>

# List of Figures

4.1	The Individual/Direct IEEE 2030.5 Model, as defined by California SIP . . . . .	7
4.2	The Aggregated Clients IEEE 2030.5 Model, as defined by California SIP . . . . .	7
5.1	A Rust trait representing the IEEE 2030.5 "Resource" data type . . . . .	13
5.2	A Rust trait representing the IEEE 2030.5 List data type . . . . .	13
5.3	Rust code required to represent the IEEE 2030.5 "Resource" data type using 'inheritance-rs' .	14
5.4	A Rust data type representing an IEEE 2030.5 List using 'inheritance-rs' . . . . .	15
5.5	A Rust data type representing an IEEE 2030.5 List as generated by <code>xsd-parser-rs</code> . . . .	16
5.6	An XML representation of an IEEE 2030.5 List data type . . . . .	16
5.7	A Rust test verifying that the IEEE 2030.5 List resource can be serialized & deserialized . . . .	19
5.8	IEEE 2030.5 function sets as grouped by QualityLogic [Allr7, ] . . . . .	22





# Chapter 1

## Introduction

In recent history, society faces a growing need for reliable, sustainable and affordable electricity. We can attempt to address this problem by way of the 'smart grid', the term used for an electric grid assisted by computers, where by communication is enabled between utilities and individuals on the network.

On this smart grid are 'end-user energy devices', devices existing in the end-user energy environment. This category of end-user energy devices encompasses what are referred to as 'Distributed Energy Resources', end-user energy devices that deliver AC power to be consumed in the residence and/or the grid, such as solar inverters, and household solar batteries. [oEE18]

Through the use of distributed energy resources, fossil fuel based energy generation can be replaced with clean, renewable energy, the need for which is of growing importance as humanity seeks to address the threat of global climate change.

However, of great importance to the success of Distributed Energy Resources is their management, and integration, into the wider electric grid. End-user energy devices require communication with electric utilities for the purpose of managing electricity supply and demand, monitoring usage, and ensuring end-users are compensated for their energy contributions, as to name a few.

The 2030.5 Protocol is an IEEE standardised communication protocol purpose built for integrating end-user energy devices, and therefore Distributed Energy Resources, into the wider electric grid. The protocol has furthermore seen unanimous adoption, or proposed adoption, by DNSPs in Australia.[Cut22]

Thus, the goal of this thesis is to implement the IEEE 2030.5 standard for use in end-user energy devices.

## Chapter 2

# Context

IEEE 2030.5 is but one protocol designed for communication between end-user energy devices and electric utilities. Although IEEE 2030.5 is a modern protocol, it is not wholly new or original, rather it is the product of historically successful protocols, designed with the same aim. In this section we'll examine the predecessors to IEEE 2030.5, and their influence on the standard.

### 2.1 Smart Energy Profile 1.x

Developed by ZigBee Alliance, and published in 2008, Smart Energy Profile 1.x was a specification for an application-layer communication protocol between end-user energy devices and electric utilities. The specification called for the usage of the "ZigBee" communication protocol, based off the IEEE 802.15.4 specification for physical layer communication. [All13, ]

The specification was adopted by utilities worldwide, including the Southern California Edison Company, who purchased usage of the system for \$400 million USD.[Heio8, ]

According to SunSpec in 2019, over 60 million smart meters are deployed under ZigBee Smart Energy 1.x, with 550 certified SEP 1.x products. [All17, ]

### 2.2 SunSpec Modbus

Referenced in the specification as the foundations for IEEE 2030.5 is the Sunspec Alliance Inverter Control Model, which encompasses the SunSpec Modbus Protocol. Sunspec Modbus is an extension of the Modbus

communication protocol designed for end-user energy devices, and was published, yet not standardised, in 2010. The protocol set out to accomplish many of the same goals as IEEE 2030.5 does today. Tom Tansy, chairman of the SunSpec Alliance pins the goal of the protocol as to create a 'common language that all distributed energy component manufacturers could use to enable communication interoperability'. [Tan21, ]

### **2.3 IEEE 2030**

IEEE 2030 was a guide, published in 2011, to help standardise smart grid communication and interoperability, and how potential solutions could be evaluated. A major goal of these potential solutions is that, by their nature of existing in the end-user environment, they were to prioritise the security of all data stored and transmitted. [oEE11, ]

### **2.4 Producing IEEE 2030.5**

In the interest of interoperability with a future standard, the SunSpec alliance donated their SunSpec Modbus protocol to form IEEE 2030.5. Simultaneously, ZigBee Alliance was looking to develop Smart Energy Profile 2.0, which would use TCP/IP. At this point the SunSpec Alliance formed a partnership with ZigBee Alliance, and IEEE 2030.5 was created as a TCP/IP communication protocol that is both SEP 2.0, and interoperable with SunSpec Modbus. [Tan21, ] Likewise, IEEE 2030.5 works to improve the security of end-user energy devices, as put forward by IEEE 2030.

With IEEE 2030.5 extending and combining these existing, widely used, protocols, we're reassured it actually solves the problems faced by utilities and smart grid device manufacturers alike, and wasn't created in a vacuum, unaware of real world requirements, or the needs of device manufacturers and electric utilities.

## Chapter 3

# Background

### 3.1 High-Level Architecture

The IEEE 2030.5 protocol is a REST API, and as such, adopts a client-server model.

Transmitted between clients and servers are 'Resources', all of which are defined in a standardised schema. Despite the client-server model, the IEEE 2030.5 specification purposefully does not make distinctions between clients and servers, as to avoid resources having differing behaviours on each. Rather, a server simply exposes resources to clients, and clients retrieve, update, create and delete resources on servers. A client communicates with, at most, one server, and a server communicates with many clients.

Being the product of existing technologies, the IEEE 2030.5 resources cover a wide range of applications, and as such, the specification logically groups resources into discrete 'function sets', of which there are twenty-five. Device manufacturers or electric utilities implementing IEEE 2030.5 need only communicate resources from function sets relevant to the purpose of the device.

The specification defines two methods by which clients retrieve resources from server. The first of which, polling, has a client contact a server, requesting updates, on a timed interval. The second of which has clients 'subscribe' to a resource, where by a server will then send a notification to all subscribed clients when that resource is updated. [oEE18]

## 3.2 Protocol Design

Resources are transmitted between clients and servers using HTTP/1.1, over TCP/IP using TLS. As a result, the protocol employs the usual HTTP request methods of GET, POST, PUT and DELETE for retrieving, updating, creating and deleting resources, respectively.

The specification requires that TLS certificates be signed, and all encryption done, using ECC cipher suites, with the ability to use RSA cipher suites as a fallback.

The standardised resource schema is defined using XSD, as all transmitted resources are represented using either XML, or EXI, with the HTTP/1.1 Content-Type header set to `sep+xml` or `sep-exi`, respectively.

In order to connect to servers, clients must be able to resolve hostnames to IP addresses using DNS. Similarly, the specification permits the ability for clients to discover servers on a local network using DNS-SD.

## 3.3 Usage

The IEEE 2030.5 function sets cover a wide range of applications and uses, aiming to support as many end-user energy devices as possible. A subset of these possible use cases are as follows:

- (Smart) Electricity Meters can use the 'Metering' function set to 'exchange commodity measurement information' using 2030.5 resources. [oEE18]
- Electric Vehicle chargers may wish to have their power usage behaviour influenced by resources belonging to the 'Demand Load and Response Control' function set.
- Solar Inverters may use the 'Distributed Energy Resources' function set such that their energy output into the wider grid can be controlled by the utility.

## Chapter 4

# Adoption

Further proving that the IEEE 2030.5 protocol is worth implementing is its adoption by electric utilities and tariffs and guidelines created by government energy bodies mandating its use. Consequently, many implementations of the protocol exist already, with the vast majority of them proprietary, or implementing proprietary extensions of the standard. In this section, we will examine both, and discuss how they may influence our open-source implementation.


### 4.1 California Public Utilities Commission

'Electric Rule 21' is a tariff put forward by CPUC. Within it are a set of requirements concerning the connection of end-user energy production and storage to the grid. In this tariff, it is explicitly clear that "The default application-level protocol shall be IEEE 2030.5, as defined in the California IEEE 2030.5 implementation guide" [GC23, ]. Given that the state of California was among the first to make these considerations to the protocol, it's likely that future writers of legislation or tariffs will be influenced by Rule 21, particularly how they have extended the protocol to achieve scale in the realm of smart inverters. For that reason, we let the implementation models for the Californian IEEE 2030.5 implementation guide influence our own development of the protocol, whilst of course still adhering to the specification.

Relating directly to use of the IEEE 2030.5 protocol at scale are the high level architecture models defined in the California SIP implementation guide.

## Individual/Direct Model

Under this model there is a direct communication between an IEEE 2030.5 compliant client, in this case a solar inverter, and a IEEE 2030.5 compliant server, hosted by the electric utility. This model alone does not impose any additional restrictions over those already existing in Rule 21. It requires the inverter to be a 2030.5 Client, and be managed individually by the server.



End-user Energy Device

Figure 4.1: The Individual/Direct IEEE 2030.5 Model, as defined by California SIP

## Aggregated Clients Model

The aggregated clients model, outlined in the implementation guide, is one preferred for use by electric utilities. Under this model, the 2030.5 client is but an aggregator communicating with multiple smart inverters, acting on their behalf. The rationale behind this model is to allow utilities to manage entire geographical areas as though it were a single entity.




Figure 4.2: The Aggregated Clients IEEE 2030.5 Model, as defined by California SIP

The IEEE 2030.5 server is not aware of this aggregation, as the chosen communication protocol between an aggregator client and an end-user energy device is unspecified and not relevant to the model, as indicated in Figure 4.2. Under this model, aggregators may be communicating with thousands of IEEE 2030.5 compliant clients. For this reason, the California SIP mandates the subscription/notification retrieval method be used by clients, rather than polling. This is mandated in order to reduce network traffic, and of course, allow for use of the protocol at scale.

Given the circumstances of this model, the aggregator IEEE 2030.5 client is likely to be hosted in the cloud, or on some form of dedicated server.

This model influences our client implementation in that it must be scalable, and take advantage of available hardware resources where possible. For instance, our client must take advantage of parallel computing on multi-threaded devices. Likewise, it must implement the subscription/notification retrieval method.

## **4.2 Australian Renewable Energy Agency**

"Common Smart Inverter Profile" (Common SIP) is an implementation guide developed by the "DER Integration API Technical Working Group" in order to "promote interoperability amongst DER and DNSPs in Australia". The implementation guide has DER adhere to the IEEE 2030.5 spec, whilst further leveraging the aforementioned CPUC California SIP, including support for use of the client aggregator model, and the mandated use of subscription/notification retrieval by those aggregator clients. [Age23, ]

Most importantly, the Australian Common SIP extends upon the `DERCapability` resource to support `Dynamic Operating Envelopes`, and it does so whilst still adhering to the IEEE 2030.5 specification. As per the specification, resource extensions are to be made under a different XML namespace, in this case `https://csipaus.org/ns`. [Age23, ]

As such, our client implementation will allow for extended resources to be created and used, at the convenience of the device manufacturer (See Section 5.3.1). However, we won't rule out the possibility of implementing the extensions, as specified by the Common SIP, as a stretch goal.



### 4.3 SunSpec Alliance

As of present, the specification behind the aforementioned SunSpec Modbus is still available and distributed, as it's "semantically identical and thus fully interoperable with IEEE 2030.5". The primary motivation for implementing SunSpec Modbus is it's compliance with IEEE 1547, the standard for the interconnection of DER into the grid. [All, ]

### 4.4 Electric Power Research Institute

One of the more immediately relevant adoptions of the protocol is the mostly compliant implementation of a client by EPRI in the United States. Released under the BSD-3 license and written in the C programming language, the implementation comes in at just under twenty-thousand lines of C code. Given that a IEEE 2030.5 client would require extension by a device manufacturer, as to integrate it with the logic of the device itself, EPRI distributed header files as the core of the client. For the purpose of demonstration and testing, they also bundled a client binary that can be built and executed, running as per user input.

The C codebase includes a program that parses the IEEE 2030.5 XSD and converts it into C data types (structs) with documentation. This is then built with the remainder of the client library.

The implementation targets the Linux operating system, however, for the sake of portability, EPRI defined a set of header file interfaces that contained Linux specific API calls, such that they could be replaced for some other operating system. These replaceable interfaces include those for networking, TCP and UDP, and event-based architecture, using the Linux `epoll` `syscall`, among others.

The IEEE 2030.5 Client implementation by EPRI states, in it's User's Manual, that it almost perfectly conforms to the IEEE 2030.5 specification according to tests written by QualityLogic. The one exception to this is that the implementation does not support the subscription/notification mechanism for resource retrieval, as of this report. This is particularly unusual given that the California SIP mandates the use of subscription/notification under the client aggregator model, a model of which this implementation was targeted for use in. For that reason, this thesis will prioritise implementing the subscription/notification mechanism, with the purpose of filling a gap in available open-source technologies. [Ins18, ]

## Chapter 5

# Implementation Design

Given the context and background surrounding the IEEE 2030.5 Protocol, we can begin examining the considerations taken when designing our implementation.

### 5.1 Programming Language

The operating context of IEEE 2030.5 greatly influences the software development tools that are suitable for implementation. When choosing a programming language, we must consider:

- **Device resource constraints** - Many existing IEEE 2030.5 compliant devices are constrained by the clockspeed of the onboard CPU, and the amount of onboard memory. For example, the "Gen5 Riva" smart meter by Itron, possesses a 600Mhz processor, and 256Mb of RAM. [oC20, ] Therefore, we desire a programming language with minimal runtime overhead, and low-cost abstractions.
- **Cross-platform capabilities** - Given that our client may be deployed on embedded hardware, writing code in a language that can be compiled to a wider range of CPU and hardware configurations is desirable. For this reason, we saw EPRI choose to implement their client in the C programming language.
- **Security** - In 2019, Microsoft attributed 70% of all CVEs in software in the last 12 years to be caused by memory safety issues. [Mil19, ] As we've previously established, improving the security of end-user energy devices is part of the core rationale behind IEEE 2030.5. For that reason, developing our client in a language that is memory safe would better align our implementation with the goals of the specification.

- **Concurrent programming support** - By nature of the protocol, an IEEE 2030.5 client is an I/O bound application. For that reason, we desire a programming language that provides us with abstractions for event-driven architecture, such that operations can be performed while waiting on I/O, even on a single threaded device.

### 5.1.1 Rust

Rust is a high-level, general purpose programming language. Rust will be used to develop our IEEE 2030.5 client implementation as it addresses all the aforementioned considerations.

- Rust provides us with high-level abstractions, but with a focus on performance, and only 'paying' (in terms of available computer resources) for what you use. For that reason, Rust has seen usage in both application and systems level programming.
- Although Rust is by no means a universal programming language like C, it has a growing presence in embedded development, and as such has a reasonable degree of portability.
- While writing a variation of Rust called 'safe' Rust, the programming language is, outside of compiler bugs, completely memory safe, accomplished without the runtime overhead of a garbage collector. When writing 'unsafe' rust, raw pointers may be dereferenced, where by the compiler is no longer able to guarantee memory safety. For that reason, our implementation will be done using purely safe Rust.
- Through it's type system, Rust also eliminates the possibility of a data race when working with multiple threads of execution. Furthermore, Rust has native support for asynchronous programming, with concurrency-runtime agnostic await and async syntax.

Rust's high-level abstractions furthermore lend themselves to the imposed reliability of the client, in that all expected errors are to be handled at compile-time. The Rust algebraic sum types 'Option' and 'Result' force programmers to handle error cases in order to use the output of a process. This is contrasted to a language like C++, where runtime exceptions are used to denote errors in the standard library, of which the language does not force the programmer to handle at compile-time.

## 5.2 Operating System

Despite the desire to write code that is portable, we will nonetheless require operating-system-specific functionality, and as such will need to target a single operating system. Then, we can provide an interface for device manufacturers to swap out should they want to use the client on a different operating system.

Of great consideration when choosing an operating system is the aforementioned 'Aggregator' model for IEEE 2030.5, where by our client would be deployed on a dedicated server, or in the cloud. In this circumstance, it's very much likely an operating system running on the Linux kernel is to be used, due to it's prevalence in server operating-systems. Whilst there exist purpose built IOT-device operating systems, they are built and optimised for low-spec low-power usage devices, and as such would not be appropriate at scale such as under this model.

Furthermore, there exist very lightweight Linux based operating systems for low-spec devices, making those Linux based operating systems the best candidate for our targeted operating system.

Here, it's reasonable to draw inspiration from the EPRI implementation of an IEEE 2030.5 client, and still define a series of interfaces that call upon OS specific APIs, such that a device manufacturer need only re-implement those interfaces when porting our software to a different OS.

## 5.3 Common Library

As discussed, the IEEE 2030.5 specification makes no distinction between clients & servers, aside from the fact that servers expose resources, and clients interact with resources. For that reason, our client implementation will also produce a common library, that would be used by one developing a IEEE 2030.5 server implementation in Rust.

### 5.3.1 Resource Data Types

The largest aspect of this common library is the internal representation of resources, the data communicated between client & server. These resources are described precisely in an XSD. Resources range from data that may be used by the electric utility, such as the current load of the device, to resource metadata, such as the rate at which a client should poll a server for a given resource, or what URI endpoint can be used to access a given resource, in the case of a Link resource. In both specification and schema, these data structures are separated into packages, and sub-packages.

Whilst the 2030.5 specification makes no mention of the object oriented programming paradigm, OOP inheritance underpins the design of all resources, including both multi-level and hierarchical inheritance. As such, for the purpose of code reuse, many base types appear in the 'common' package; data structures extended by many others.

## Representing Resources in Rust

Rust, despite being influenced by OOP, does not possess the notion of a class, like in languages like C++ or Java, and as such does not define shared behaviour of types through inheritance from a common parent type. Rather, Rust defines shared behaviour through traits, where shared behaviour refers solely to the methods that we can call on that type, as per the traits a type implements. In this sense, Rust does not concern itself with what a type is or what that type stores, it concerns itself only with the traits a type possesses.

Traits themselves do support inheritance. One can have traits that require other traits to be implemented for a given type. However, this does not change the fact that traits only represent behaviour. There is no way to have a data structure inherit the internal members of another.

```
pub trait Resource {  
    fn get_href(&self) -> Option<&str>;  
}
```

Figure 5.1: A Rust trait representing the IEEE 2030.5 "Resource" data type

```
pub trait List : Resource {  
    fn all(&self) -> UInt32;  
    fn results(&self) -> UInt32;  
}
```

Figure 5.2: A Rust trait representing the IEEE 2030.5 List data type

In Figure 5.2, we have a Rust trait that describes the behaviour of the List base type. All lists are resources, and thus we have a trait restriction that all types implementing List must first implement the Resource trait.

This is the extent of native inheritance in rust. We can specify the exact behaviour of types that belong to a trait in detail, but we cannot influence how that behaviour is achieved.

## Emulating inheritance in Rust

As a result, we're forced to emulate the inheritance of data structure members in Rust, of which there are two approaches:

- Composite an instance of the base type into type definitions
- Repeat all inherited members in type definitions

Regardless of the approach, we still do not have have polymorphism using the base-types of resources. To allow for polymorphism a trait must be defined for each base-type, and then those trait functions need be implemented for every type that extends that base type. This is unavoidable duplicate code, not ignoring the possibility of generating said code. The extent to which polymorphism is required in the final client implementation is yet to be determined. In this case, we can refer to the EPRI client implementation, where they provide device manufacturers an interface for handling resources using polymorphism. Thus, when determining how we will emulate inheritance, we let the ease at which polymorphism can be allowed influence our decision

## Inheritance via Composition

If we were to implement the first of the two approaches, we can make use of an existing Rust library to reduce the amount of boilerplate required to implement polymorphism. This library operates on the basis that inheritance can be replicated via composition. If a data type were to contain a member that implements a given trait, there is no reason for that outer struct to not be able to implement that trait by simply calling upon the underlying member.

```
#[inheritable]
pub trait Resource {
    fn get_href(&self) -> Option<&str>;
}

pub struct ResourceObj {
    href: Option<String>
}

impl Resource for ResourceObj {
    fn href(&self) -> Option<&str> {
        self.href.as_str()
    }
}
```

Figure 5.3: Rust code required to represent the IEEE 2030.5 "Resource" data type using 'inheritance-rs'

```

#[derive(Inheritance)]
pub struct List {
    #[inherits(Resource)]
    res: ResourceObj,
    all: UInt32,
    results: UInt32,
}

```

Figure 5.4: A Rust data type representing an IEEE 2030.5 List using 'inheritance-rs'

Figures 5.3 and 5.4 show how this library, 'inheritance-rs' [HM19, ] is used to reduce the boilerplate necessary to inherit data members. In Figure 5.3, we mark the Resource trait as 'inheritable' and then implement that trait on a type that holds the necessary members, our bare minimum 'base' type. In Figure 5.3, we compose an instance of that base type into a type that would normally inherit from it. Then, we tell the library to generate the code, at compile time, that would allow List to implement the Resource trait. This generated code simply calls the underlying ResourceObj member when the href function would be called on a list.

The major flaw in this approach is that for every single type that is used as a base type, a trait, and a base implementation of that trait needs to be written. Given that there are just under 700 data types in the IEEE 2030.5 specification schema, we must consider alternatives.

## XSD to Rust types

If we were to implement the second approach, an existing Rust library can be used to automate the process of defining data types altogether. This of course draws on the fact that the IEEE 2030.5 XSD is entirely self-contained, and follows XSD guidelines by W3C. As such, generating rust data types from it is a reasonable approach. One such way to automate this process would be to design and implement our own XSD parser, however, we are not the first to require this tool.

On the Rust public crates registry there are several XSD parsers, many of which existing to solve very similar problems; implementing standardised communication protocols in Rust. However, for that reason, many of these implementations are developed until they meet the creators needs, at which point the tool is no longer maintained.

Of the most complete parsers, one particular implementation stands out. This particular implementation supports hierarchial inheritance and makes reasonable assumptions on the internal representations of primitive data

types. `xsd-parser-rs` by Lumeo, was created for use in their Open Network Video Interface Forum Client, software with requirements not dissimilar from that of IEEE 2030.5. [Lum22, ]

```
#[derive(Default, PartialEq, Debug, YaSerialize, YaDeserialize)]
#[yaserde(namespace = "urn:ieee:std:2030.5:ns")]
pub struct List {
    // The number specifying "all" of the items in the list.
    // Required on a response to a GET, ignored otherwise.
    #[yaserde(attribute, rename = "all")]
    pub all: Uint32,

    // Indicates the number of items in this page of results.
    #[yaserde(attribute, rename = "results")]
    pub results: Uint32,

    // A reference to the resource address (URI).
    // Required in a response to a GET, ignored otherwise.
    #[yaserde(attribute, rename = "href")]
    pub href: Option<String>,
}
```

Figure 5.5: A Rust data type representing an IEEE 2030.5 List as generated by `xsd-parser-rs`

Figure 5.5 is what `xsd-parser-rs` currently produces for the List resource. In this figure, it's parsed that the List type inherits from the Resource type, and included the href data member accordingly. It's also included the documentation as found in the schema.

Compared to true OOP inheritance, this has types include their parent type data members in their own type definitions. Despite this, the type definitions are far more readable, and align more closely with the output of the client, in that inheritance is flattened out.

```
<List xmlns="urn:ieee:std:2030.5:ns" all="0" results="0"
href="/sample/list/uri" />
```

Figure 5.6: An XML representation of an IEEE 2030.5 List data type

Figure 5.6 shows the List data type were it a concrete data type serialized into XML.

Further advocating for the use of an XSD parser is the fact that the 2030.5 XSD is updated more often than revisions of the specification itself, even if those updates are relatively infrequent. Using `xsd-parser-rs` will allow us to better maintain the client, by way of simply running updated schemas through the parser. Furthermore, as DER implementation guides choose to extend the available IEEE 2030.5 resources under new namespaces,



our `xsd-parser-rs` fork will be made available such that the XSD for those extensions be provided as input, and the corresponding Rust data types be generated.

Regrettably, `xsd-parser-rs` does not perfectly fit our needs:

- It currently does not auto-generate Rust enums, since types where a rust enum would be best suited simply have their variants as comments in the XSD itself.
- It currently doesn't create and implement traits for us using base types, therefore, no polymorphism.

For that reason, as part of our client implementation, we'll be maintaining a git fork of the tool and will be implementing this functionality as required.

### 5.3.2 Resource Serialisation & Deserialisation

As we've established, resources are sent as their XML, or EXI, representations over HTTP. This means our common library requires the ability to serialize and deserialize resources to and from the appropriate XML or EXI. For XML, this means we're sending HTTP requests with `Content-Type` set to `application/sep+xml`, and for EXI this means we're setting it to `application/sep-exi`,

#### EXI

As of present, there exists no Rust library for producing EXI from XML or from Rust data types, or vice-versa. For that reason, one of the stretch goals of this thesis will be to develop the first EXI library for Rust, and use that in the common library. This EXI library could be developed from scratch, or could call C code from Rust; using the `Embeddable EXI Processor` implementation would be a sensible approach.

#### XML

On the public Rust crates registry, there exists a popular XML serialisation and serialization library purpose built for use in embedded communication protocols called `YaSerde` [Lum23, ]. Despite being developed by different teams, the syntax to use `YaSerde` on Rust data types is auto-generated by `xsd-parser-rs`, and as such, is to be used in our client implementation.

Referring again to Figure 5.5, we see `xsd-parser-rs` has qualified our struct with the appropriate namespace, such that YaSerde serializations will include it, and deserializations will expect it. Furthermore, names of data members have been specified as attributes, rather than child elements in the resulting XML; determined by the parser as per the XSD. Figure 5.6 is an example of what the List type looks like when serialized to XML.

Once again, YaSerde does not perfectly fit our needs. Currently, the library serializes enums as a string representation of their variant. IEEE 2030.5 requires that the underlying integer representation be used, as the variants are but annotations in the schema.

Rather than go and manually implement XML serialization/deserialization for all enums, we can simply maintain a fork of YaSerde, and modify the library's behaviour there.

### 5.3.3 Security

The security-focused portion of the common library includes certificate management, and an encryption interface for use in client networking. The specification requires that ECC, as defined in IETF RFC 7251, specifically elliptic curve `secp256r1`, be used to sign certificates and encrypt traffic over TLS. However, the specification does permit RSA-based cipher suites to be used as a fallback, for the purpose of backwards compatibility, on devices where ECC is not supported.

For that reason, our common library will implement a security interface to be used throughout the client and server, such that it is comprehensive and can be replaced for portability.

### 5.3.4 Network

The network-focused portion of the common library includes creating and sending HTTP requests over TCP, whilst using the previously defined security interface for TLS.

It's worth noting that the extent to which a server implementation would utilise this interface is yet to be determined.

### 5.3.5 Testing

The goal for this thesis is to produce a IEEE 2030.5 client that is well-tested for correctness, and conformance to the IEEE 2030.5 specification. As such, our common library will be developed alongside a suite of tests, written

with the goal to maximise code coverage, measured by lines of code. All testing for the Common Library will be done via unit testing.

## Testing YaSerde

As expected, by the nature of serialisation and deserialisation, we have no compile time guarantees either will succeed for a given data type. Therefore, our first test suite for the common library is to test that all data types can be serialized & deserialized successfully, and that it can be done without any data loss.

```
#[test]
fn list_serde() {
    let orig = List::default();
    let new: List = from_str(&to_string_with_config(&orig, &
        YASERDE_CFG).unwrap()).unwrap();
    assert_eq!(orig, new)
}
```

Figure 5.7: A Rust test verifying that the IEEE 2030.5 List resource can be serialized & deserialized

The logic for a test of this type can be succinctly expressed in Rust, as seen in Figure 5.7. This test takes advantage of the fact that all our resources are able to automatically implement the 'default' trait, and as such an instance of any resource can be instantiated with default values. In Figure 5.7, we first perform this default instantiation, convert the resource to an XML string, and then convert the XML string back into our internal List representation. Finally, we check if the newly deserialized resource is equal to it's default.

For each resource, the only part of this test that differs is the name of the type, making this test suite a very reasonable automation. Not only will this inform of us any bugs or issues in YaSerde that lead to data loss, it will also inform us of any limitations of YaSerde, in terms of things it can and cannot serialize/deserialize. Furthermore, as we make changes to our YaSerde fork, this test suite will inform us of any regressions.

However, this test suite is not capable of testing for resource serialization correctness; whether or not the outputted XML conforms to the specification.

## Unit Tests

The remainder of the common library will be tested using manually written unit tests. A set of these unit tests will test whether resource serialization matches that of the specification using the provided example outputs, whilst others will ensure the network and security interfaces are correct. These unit tests are paramount when

working with our Rust crate dependencies, as they can be used to test for regressions if and when we choose to update them.

## 5.4 Client Library

Despite the claim that we are developing an IEEE 2030.5 client, in reality, the goal of this thesis is to produce a client library, such that manufacturers of IEEE 2030.5 compliant devices will extend upon our client. These device manufacturers would simply, in Rust or otherwise, call upon the documented interface of the client library to communicate with a IEEE 2030.5 server.

The exact details of this client library interface are yet to be determined, but what a developer should be able to achieve with that interface is fairly straight forward. They'll require:

- A mechanism for discovering or connecting to a specific server.
- The ability to retrieve resources from a specific URI, and have it deserialized into a local data type.
- The ability to create or update a resource, and have it included as a PUT or POST request to a connected server.
- A mechanism to automate this process, using event-driven architecture, to handle timer events, such as for polling, or network events for resource subscription.

### 5.4.1 Server Discovery & Connection

The IEEE 2030.5 specification states that a connection to a server by a client can be established by specifying a specific IP address or hostname and port. In the event this is not provided, the specification states that DNS-SD can be used to query a local network for servers, whilst providing the ability to only query for servers advertising support for specific function sets. As of this report, there is little perceived value in this functionality, as electric utilities will almost certainly be capable or willing to supply the address and port of a server. The EPRI IEEE 2030.5 client manual indicates the same perception.[Ins18, ] As such, the implementation of this functionality will be left as a thesis stretch goal. Nonetheless, it will need to be implemented for conformance to IEEE 2030.5.

Connecting to a server is, otherwise, as expected. We use the cipher suite interface and the TCP interface, as built in the common library, to connect to the server using TCP/TLS. This of course requires certificates signed by a

common library cipher suite. This certificate exchange will determine the parameters by which the cipher suite will be used to encrypt traffic.

#### **5.4.2 Event-driven Architecture**

With the IEEE 2030.5 Client being an I/O bound application, we are to implement concurrency using an event-driven architecture, such that many operations can be performed concurrently on a single threaded machine when waiting on I/O.

Although the client does not connect to multiple servers, we require event-based architecture to handle:

- Input from the client device, such as the creating or updating of resources.
- Timed events, such as to poll for a resource update on a given schedule.
- Network events, such as an updated resource being pushed to the client via the subscription/notification mechanism, or receiving the response from a sent HTTP request.

An event-based architecture furthermore enables us to take advantage of multiple OS threads, and therefore multiple CPU cores, as to best accommodate the scale required by the aforementioned client aggregator model. As mentioned previously, Rust supports asynchronous programming, with runtime-agnostic support for `await` and `async` syntax, of which there are multiple Rust `async` runtimes available.

The specifics of this asynchronous programming architecture are yet to be learned, and as such we can look to the EPRI client implementation, where they've chosen to use the Linux native `epoll` `syscall`. Fortunately, `epoll` is the driver behind the two most popular `async` runtimes for Rust, `Tokio` and `std-async`. As such, we'll need to determine which is most suitable for our client.

#### **5.4.3 Function Sets**

As established, function sets are logical groupings of IEEE 2030.5 resources by their functionality, of which there are twenty-five. A subset of these function sets are for general or 'Common' use, such as the time function set, for synchronizing a device clock, whilst others have a very specific scope, such as the Electric Vehicle function set.

QualityLogic, a quality assurance company specialising in testing standardised software, has grouped these function sets further than the specification itself.

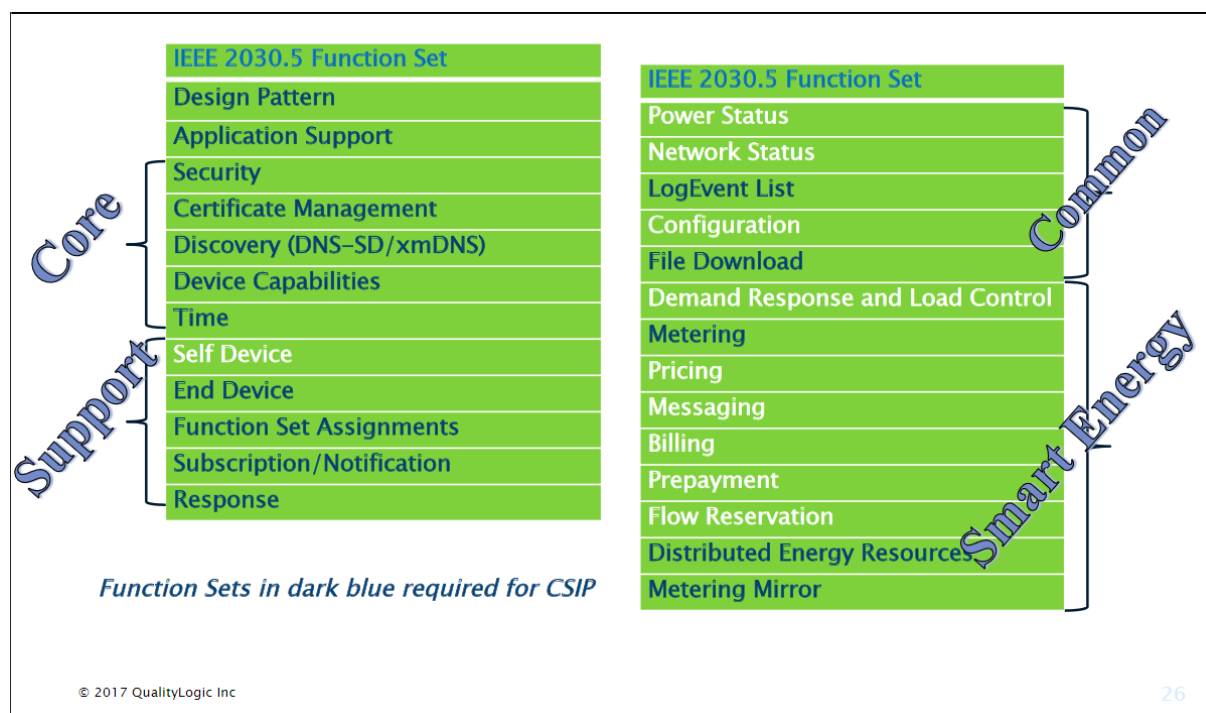


Figure 5.8: IEEE 2030.5 function sets as grouped by QualityLogic [All17, ]

Figure 5.8 shows these groupings, and which of these function sets are required for California SIP. These groupings provide us with a starting point for determining the priority of function sets, where by those required for California SIP will be implemented first. An exception to this is the Demand Response and Load Control function set, which has been deemed relevant for DER in Australia by electric utilities in Queensland, and as such it's implementation will be prioritised. [EE20, ]

Regardless, the groupings in Figure 5.8 further indicate the dependencies between function sets. For example the Device Capabilities function set describes the function sets available on a server, and is therefore used to discover the URI of resources for other function sets.

The conclusion to draw from this is that function sets on the left side of Figure 5.8 are very much depended upon by function sets on the right side, as they describe the core functionality on the client. Our development will begin on the Security, Certificate Management, and Application Support function sets as part of the common library, where the Application Support function set refers to the use of prescribed technologies, such as TCP, HTTP/1.1 and XML.

#### **5.4.4 Testing**

As expected, our client library will be developed alongside a suite of tests. Unlike the common library, these tests will be comprised of both unit and system tests. System tests will primarily be used to check for conformance to the IEEE 2030.5 specification, using the example system outputs in the specification. The actual software framework to be used for conducting system tests is yet to be determined, but the requirements to use such a framework are relatively straightforward, obviously requiring both a test client binary, and a mock server.

##### **Test Client Binary**

In order to perform system tests we will need to wrap our client library in code that can be used to produce a binary that matches a potential use case. This binary client will be a demonstration of all implemented functionality. Furthermore, this client binary will act as documentation for device manufacturers using our library. Producing a test client binary for demonstration was an approach also taken by EPRI in their IEEE 2030.5 client implementation.

##### **Mock Server**

When it comes to writing system tests, our testing philosophy is hindered by a lack of open-source IEEE 2030.5 implementations, in-development or otherwise. As such, we'll need to create our own mock IEEE 2030.5 server to test both the sending and receiving of resources from our client. Neel Bhaskar, as part of UNSW CSE, developed a portion of a working IEEE 2030.5 Server in 2022, and with changes to support hardcoded resources and routes, will be used as the starting point for this mock server. We will also need to make changes to this server in order for it to use our common library, where we will therefore gain a better idea of what functionality is common to both client & server.

Specific client functionality, such as the subscription/notification method of resource retrieval, may require extensive work on this mock server. As a result, we won't rule out the possibility of writing code that could be reused in a future server implementation.

# Chapter 6

## Timeline

### 6.1 Thesis A

Estimated: 150 hours.

#### February 13th to April 26th, 2023

- Background Research
- Common Library Resource Data Structures
- Thesis A Seminar
- Thesis A Report

#### April 26th to May 29th, 2023

- Common Library: Resource Serialization & Deserialization Test Suite
- Common Library: Cipher Suite Interface + Unit Tests (Security FS)
- Common Library: Network Interface + Unit Tests (Application Support FS)
- Mock IEEE 2030.5 Server & System Test Framework
- Test Client Binary



- Client Library: Event-Based Architecture

## **6.2 Thesis B**

Estimated: 150-200 Hours

**May 29th to September 11th, 2023**

- Client Library
  - Time FS: Scheduled Events, Server polling
  - End Device FS: Client & Server Relationship
  - Subscription/Notification Resource retrieval
  - Function Sets: Metering, LogEvent, FileDownload, Demand Response and Load Control, Distributed Energy Resources
  - Corresponding Unit Tests
- Mock Server Updates
- Test Client Updates
- Thesis B Demonstration of all implemented functionality

## **6.3 Thesis C**

Estimated: 150-200 Hours

**September 11th to November 24th, 2023**

- Client Library: Function Sets
  - Metering Mirror
  - Power + Network Status
  - Configuration

- Demand Response and Load Control
  - Pricing
  - Messaging
  - Billing
  - Prepayment
  - Flow Reservation
- Mock Server Updates
- Test Client Updates
- Thesis C Seminar
- Thesis C Report
- Stretch Goals
  - EXI library for Rust
  - DNS-SD Server Discovery
  - Australian Common SIP Extensions

# Bibliography

- [Age23] Australian Renewable Energy Agency. Common smart inverter profile - australia. <https://arena.gov.au/assets/2021/09/common-smart-inverter-profile-australia.pdf>, January 2023.
- [All] SunSpec Alliance. Sunspec modbus. <https://sunspec.org/sunspec-modbus-specifications/>, 2021.
- [All13] ZigBee Alliance. Zigbee specification faq. <https://web.archive.org/web/20130627172453/http://www.zigbee.org/Specifications/ZigBee/FAQ.aspx>, June 2013.
- [All17] SunSpec Alliance. Ieee 2030.5/ca rule 21 foundational workshop. <https://sunspec.org/wp-content/uploads/2019/08/IEEE2030.5workshop.pdf>, June 2017.
- [Cut22] CutlerMerz. Review of dynamic operating envelope adoption by dnsps. <https://arena.gov.au/assets/2022/07/review-of-dynamic-operating-envelopes-from-dnsps.pdf>, August 2022.
- [EE20] Ergon Energy and Energex. Enabling dynamic customer connections for der: Consultation paper. <https://www.talkingenergy.com.au/dynamicder>, December 2020.
- [GC23] Pacific Gas and Electric Company. Electric rule no.21. [https://www.pge.com/tariffs/assets/pdf/tariffbook/ELEC\\_RULES\\_21.pdf](https://www.pge.com/tariffs/assets/pdf/tariffbook/ELEC_RULES_21.pdf), February 2023.
- [Heio8] Bob Heile. Zigbee smart energy. <https://www.altenergymag.com/article/2008/10/zigbee-smart-energy/468/>, January 2008.
- [HM19] Daniel Henry-Mantilla. inheritance-rs. <https://github.com/danielhenrymantilla/inheritance-rs>, October 2019.
- [Insi8] Electric Power Research Institute. Electric power research institute ieee 2030.5 client user's manual. <https://www.epri.com/research/products/000000003002014087>, July 2018.
- [Lum22] Lumeo. xsd-parser-rs. <https://github.com/lumeohq/xsd-parser-rs>, August 2022.
- [Lum23] Lum::Invent. Yaserde. <https://github.com/media-io/yaserde>, February 2023.
- [Mil19] Matt Miller. A proactive approach to more secure code. <https://msrc.microsoft.com/blog/2019/07/a-proactive-approach-to-more-secure-code/>, July 2019.
- [oC20] Minnesota Department of Commerce. Methods for Performance Evaluations, Metrics, and Consumer Protections for AMI and FAN, December 2020.

- [oEE11] Institute of Electrical and Electronics Engineers. Ieee guide for smart grid interoperability of energy technology and information technology operation with the electric power system (eps), end-use applications, and loads, 2011.
- [oEE18] Institute of Electrical and Electronics Engineers. Smart Energy Profile Application Protocol (2030.5-2018), December 2018.
- [Tan21] Tom Tansy. Get on the modbus: An explanation of ieee 1547 and why it matters for solar installers. <https://solarbuildermag.com/news/get-on-the-modbus-an-explanation-of-ieee-1547-and-why-it-matters-for-solar-installers> March 2021.